

# Introduction aux Systèmes d'Exploitation

## Unit 6: Basics of Assembly Language

François Taïani



# Outline

- Assembly instructions
  - Encoding
  - Types of operations
  - Operands & Registers
- Memory
  - Little and Big Endian encoding
  - Signed integers
- Overview of the x86 ISA
- Anatomy of a nasm assembly file

# Reminder

- Executable machine code = sequence of instructions
  - (+ other very important metadata: symbols, libs, ...)
- An instruction = sequence of bits
  - E.g. '01001000 10000011 11101100 00001000'
  - Or '48 83 ec 08' in hexadecimal
  - Means "subtract 8 from 64bits stack pointer register"  
(on x86-64 architecture)
- Hard to manipulate directly → textual representation
  - sub rsp,0x8
  - "assembly code"

# Example: helloworld.asm

```
SECTION      .data
message:     db  "Hello, World!", 10          ; note the newline at the end
msgLen:      equ $-message

SECTION      .text
GLOBAL       _start

_start:
    mov      rax, 1                  ; system call for write
    mov      rdi, 1                  ; file handle 1 is stdout
    mov      rsi, message           ; address of string to output
    mov      rdx, msgLen            ; number of bytes
    syscall                   ; invoke operating system to do the write
    mov      rax, 60                ; system call for exit
    mov      rdi, 0                  ; exit code 0, equiv to xor rdi, rdi
    syscall                   ; invoke operating system to exit
```

```
$ nasm -felf64 helloworld.asm
$ ld helloworld.o
$ ./a.out
Hello, World!
$
```

# Compiling with nasm

## ■ Compiling to machine code

- `nasm -felf64 <filename>.asm`
- Produces a new file `<filename>.o` (object file)
- `-felf64` needed to produce x86-64 code for Linux

## ■ Linking (creating executable from object file)

- `ld <filename>.o -o <executable>`
- `-o` optional, executable called `a.out` by default
- Complex step when shared libs involved

## ■ Executing

- `<executable>`

# Anatomy of an assembly file

- Organized in sections
  - “zones” of executable files, and then of memory
  - Typically two: “.data” for data, and “.text” for code

# Data section

SECTION .data

- Reserves space in memory for initialized data
- And (optionally) defines labels pointing to this space
- General format

[label:]<sub>opt</sub> <directive> value[,value,...]

- Directive: typically `db` (define byte) or `dw` (define word)

db	0x55	; just the byte 0x55
db	0x55,0x56,0x57	; three bytes in succession
db	'a',0x55	; character constants are OK
db	'hello',13,10,'\$'	; so are string constants
dw	0x1234	; 0x34 0x12

# Data section (cont.)

- When a label is defined, can be used as an address

- Example : reserving & initializing 14 bytes

```
message: db "Hello, World!", 10
```

- In the text section, using the label to refer to this zone

```
mov BYTE [message],0x20 ; overwriting 1st byte
```

```
mov rsi, message ; memory address of 1st byte to rsi
```

- Constant can also be declared with equ

- Can use expression, but must evaluate to a constant

```
my_const: equ 10 ; my_const replaced by 10 in code
```

```
msgLen: equ $-message ; $ = current address
```

# Text Section

SECTION .text

- Must contain **GLOBAL \_start** somewhere
  - Sets the label `_start` as external
  - Needed by the linker (ld) and Linux to launch executable
- Entry point of the program must be labelled by `_start`
  - Does not need to be at beginning of code

# Assembly Instructions

## ■ General format

[label:]<sub>opt</sub> <operation> <operand1>, <operand2> [ ; com]<sub>opt</sub>

## ■ JVM executable

→ sub : instruction

→ rsp : operand1 (64 bit version of stack pointer "sp")

→ 0x8 : operand2 (number 8 in hexadecimal)

## ■ Intel notation (used in this module)

→ Result stored in the **first** operand

→ (AT&T notation, used in gcc, the reverse)



# Encoding Instruction

- Assembly needs to get encoded into machine code
- In x86 size of instruction code ("opcode") may vary.
- Encoding more complex than just concatenation
  - E.g. 48 in earlier example means "op on 64bit register"
  - 83 means one of ADD/OR/ADC/SBB/AND/SUB/XOR/CMP
  - Actual operation + reg selected by further bits in "ec"
- Thankfully, the assembler does this for you ☺
  - Plus other things (memory layout, constants, labels, ...)

# Types of Operations

- 1<sup>st</sup> Group : **moving/setting data** (e.g. `mov rax,8`)
  - Setting register to constant values
  - Moving data in and out into the processor
  - Mainly in/out of memory, but also HW devices
- 2<sup>nd</sup> Group : **integer arithmetic** (e.g. `add rax,8`)
  - Addition, subtraction, multiplication, division
  - Boolean operations (bit-wise): and, or, not, ...
- 3<sup>rd</sup> Group : **control flow** (e.g. `jmp 8, syscall`)
  - Jumps (possibly conditional), calls, return, interrupts
- 4<sup>th</sup> Group : **complex operations (CISC proc)**
  - Everything else (in particular floating point)

# Operands : Registers

- Nb of Registers / role vary between processors
- Main categories
  - **Data** registers (x86-64 : rax, rbx, rcx, rdx, ...)
  - **Address** registers (x86-64 : rsp, rbp, rsi, rdi)
  - **Status** register (used for carry, comparison, errors)
  - (+ Instruction Pointer<sup>(\*)</sup> & current instruction reg)

<sup>(\*)</sup>Also called "Program Counter" (PC)

# Operands : Registers (cont.)

- **64 bits** registers denoted by '**R**'
  - `rax, rbx, ..., rsp`
- **32bit** registers denoted by '**E**'
  - `eax, ebx, ..., esp`
- **16 bit** registers have **no prefix** and contain an '**X**'
  - `ax, bx, cx, dx`
- Each 16-bit register contains 2 8bit registers
  - e.g. for `ax` : '**al**' (low) + '**ah**' (high) (same for `bx, cx, dx`)
- **Not** independent
  - A bit like matryoshka dolls
  - `al` part of `ax`, which is part of `eax`, which is part of `rax`




# Operands: Memory

- At machine code level : memory = 1D array of bytes
  - address = position of data in this array
  - first position = zero
  - In Intel notation: "[200]" content of address 200
- Before virtual memory (8086, etc.)
  - Direct access to physical memory
- Recent processors: virtualisation (more in OS course)
  - Each executing program = illusion to be alone
  - Transparent translation (in HW+OS) between virtual/phys



# Operands: Memory (cont)

- Most instructions manipulate several bytes
  - e.g. `mov [100], ax` modifies [100] and [101]
- Beware : **Intel processor use "little endian" convention**
  - Least significant byte (`al`, "little end") starts encoding
  - Least significant byte at smallest address




# Endianness

- Exercise: consider x86 program
  - `mov eax, 0x0a0b0c0d`
  - `mov [a], eax ; a is some constant in assembly program`
- Content of memory ?



# Endianness

- Exercise: consider x86 program
  - `mov eax, 0x0a0b0c0d`
  - `mov [a], eax` ; a is some constant in assembly program
- Content of memory ?
  - What if proc were using Big-endian convention ?



Source : <http://en.wikipedia.org/wiki/Endianness>



# Signed Numbers

- Sometimes necessary to work with numbers  $< 0$
- Solution on  $n$  bits:
  - use raw values  $\geq 2^{n-1}$  for negative numbers
  - $[-2^{n-1}, 2^{n-1}-1]$
  - A form of encoding!
- Benefit
  - `add` & `sub` operations work the same
- Danger
  - Multiplication & division are  $\neq$ : `mul` & `imul`, `div` & `idiv`
  - Misinterpreting raw values  $\geq 2^{n-1}-1$
  - Important to remember which type has been used where

# Signed Numbers (cont)

- Example: on 8 bits
  - Encoding of -2 ?
  - Encoding of 2 ?
  - What happens if added as unsigned integers ?



# Limitation on Operands

- **Fundamental limitation** of processor instructions
  - Impossible to work on 2 memory locations in same op
  - `mov [200], [100]` is **impossible**
  - On some proc (RISC), even ADD AX, [100] impossible
  - Registers (almost) always involved
  
- Using **registers smartly** is essential for performance
  - Very fast, but small number
  - Work of compilers (and gurus for HPC libraries)
  - Increasingly complex heuristics (-O flag in gcc)

# x86-64 Instruction Overview

- Syntax used
  - R: register (rax, rbx, etc.)
  - M: memory (typically denoted by a label: `message`)
  - V: constant (often computed by nasm: `msgLen`)
  - R/M: register or memory
- Note
  - nasm usually able to guess the size of operands

```
mov [message],al ; AL=8bits=1Byte
```
  - but not always possible, in particular with constant

```
mov BYTE [message],0x20
mov [message],BYTE 0x20 ; equivalent
```

# Transferring Data

## ■ Instructions de transfert de données :

- `mov R, R/M`
- `mov R/M, R`
- `mov R/M, V`

## ■ Remarque :

- `mov M, M` → impossible

## ■ Example :

- `mov ax, [100]`
- `mov [message], ax`

# Arithmetic

- Sign change
  - **neg** R/M
- Adding
  - **add** R, R/M
  - **add** R/M, R/V
- Subtracting
  - **sub** R, R/M
  - **sub** R/M, R/V
- Increment, decrement
  - **inc** R
  - **dec** R

# Logic

- Beware: bitwise operations !
- 1-complement
  - `not R/M16`
- Bit-wise AND
  - `and R16,R/M16`
  - `and R/M16,R16`
- And similarly: bitwise `or`, `xor`
- Note 
  - `xor` used by compilers to set to 0 as faster than `mov`
  - E.g. `xor rdi,rdi` faster than `mov rdi,0`

# Jumps

- Two types of jumps
  - Unconditional jumps (always jumps)
  - Conditional jumps (jumps based on status register)

## Unconditional jump

- `jmp <address>`
- Notes
  - Address typically encoded as a label  
`jmp loop ; loop defined somewhere else`
  - Actual machine code (opcode) uses an offset

# Conditional Jumps

## ■ Based on **Status Register**

- Never manipulated directly
- Contains “flags” (status bit)
- Flags set by last “computed” result
  - S : 1 if negative
  - Z : 1 if zero
  - C : 1 if carry occurred (relevant for unsigned arithmetic)
  - O : 1 if overflow occurred (relevant for signed arithmetic)



## ■ Special instruction to perform “virtual” subtraction

- **cmp** R,R/M
- **cmp** R/M, R
- **cmp** R/M, V

Result lost,  
but sets the status register

# Conditional Jumps

- **je** : jumps if equal
  - Synonym of **jz** : jumps if zero
- **jne** : jump if  $\neq$ 
  - Synonym of **jnz** : jumps if non-zero
- **jg** : jump if  $>$
- **jge** : jump if  $\geq$
- **jl** : jump if  $<$
- **jle** : jump if  $\leq$

# Summary

At the end of this session you should be able to

- manipulate the most common x86 assembly instructions;
- do simple signed integer arithmetic;
- be able to manipulate numbers encoded in big and little endian;
- describe and explain the structure of a simple nasm assembly file.