

Introduction aux Systèmes d'Exploitation

Unit 8: Invocation and Stack management

François Taïani



Invocation Mechanism

- Well known: procedures (C), methods (Java, C++)
- Goal: write once, call from everywhere
- Challenge: how to remember where to return?

Example

```
#include <stdio.h>

void foo() {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar() {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

Represent the program's structure as a call graph



... And the program's execution as a call tree

Invocation Mechanism

- Challenge: how to remember where to return?
 - **foo** returns into **main** in 1st invocation
 - but into **bar** for 2nd invocation
- Solution:
 - Use a stack!



The Call Stack

- Special zone in process' memory ("stack segment")
 - LIFO principle
 - Grows downwards (on x86): from high to low addresses
(most common, but reverse possible, cf. ARM)

- Role of the stack
 - **Remember** where to return to after invocation
 - Pass **parameters** (more on this, registers can be used too)
 - Store **local variables** (more on this)
 - Retrieve **returned values** (more on this)

Example

```
#include <stdio.h>


void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses




Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses



Example

```
#include <stdio.h>


void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses




Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses




Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```




Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv ) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```

(info on
call to
main)

Higher addresses



Example

```
#include <stdio.h>

void foo(void) {
    printf("Executing foo\n");
    printf("Returning from foo\n");
} // End foo

void bar(void) {
    printf("Executing bar\n");
    printf("Calling foo from bar\n");
    foo();
    printf("Returning from bar\n");
} // End bar

int main( int argc, char** argv) {
    printf("Calling foo from main\n");
    foo();
    printf("Calling bar from main\n");
    bar();
    printf("Returning from main\n");
} // EndMain
```


(info on
call to
main)

Higher addresses




Stack in x86 Assembly

- X86-64: direct hardware support for **stack-based calls**
 - case of most high-level processors
- Special register: **rsp**
 - Points to top of the stack
- Two operations
 - **push R/M/V**, for instance **push rax**
 - pushes operand (rax here) onto the stack onto the stack
 - $\text{rsp} \leftarrow \text{rsp} - 8$ (moves to lower addresses)



Stack in x86-64 Assembly

- X86-64: direct hardware support for **stack-based calls**
 - case of most high-level processors
- Special register: **rsp**
 - Points to top of the stack
- Two operations
 - **push R/M/V**, for instance **push rax**
 - pushes operand (rax here) onto the stack onto the stack
 - $\text{rsp} \leftarrow \text{rsp} - 8$ (moves to lower addresses)
 - **pop R/M**, for instance **pop rax**
 - pops **8 bytes** at top of stack, move them to operand
 - $\text{rsp} \leftarrow \text{rsp} + 8$ (moves to higher addresses)



rax

Calls in x86-64

Two special op to call functions

- **call some_address**

- **call some_address**
 - Pushes address of **next instruction** on stack
 - (→ rsp moves towards **lower addresses**, stack grows)
 - Jumps to **some_address**

- **ret**

- **ret**
 - Retrieves (pops) **return address** from stack
 - (→ moves rsp towards **higher addresses**, stack shrinks)
 - Jumps to **return address**

Example: calling foo


```
call foo
...
foo:
    mov     rax, 1      ; system call for write
    mov     rdi, 1      ; file handle 1 is stdout
    mov     rsi, message ; address of string to output
    mov     rdx, msgLen  ; number of bytes
    syscall            ; invoke operating system to do the write
    ret
```

- The code after `foo:` is executed every time `call foo` is executed
- Note how `foo` is a label
 - Replaced by an address (an offset in the code segment)

Potential Problem

```
start:  
foo:  
    mov    rax, 1      ; system call for write  
    mov    rdi, 1      ; file handle 1 is stdout  
    mov    rsi, message ; address of string to output  
    mov    rdx, msgLen  ; number of bytes  
    syscall          ; invoke operating system to do the write  
    ret  
  
call foo  
    mov    rax, 60      ; system call for exit  
    mov    rdi, 0      ; exit code 0  
    syscall          ; invoke operating system to exit  
  
GLOBAL _start
```

- Will the above code work? Why?




And a solution

- Either implement foo after the exit to OS or ...

```
foo:  
    mov      rax, 1      ; system call for write  
    mov      rdi, 1      ; file handle 1 is stdout  
    mov      rsi, message ; address of string to output  
    mov      rdx, msgLen  ; number of bytes  
    syscall            ; invoke OS to do the write  
    ret  
  
entry point → _start:  
    call   foo  
    mov      rax, 60      ; system call for exit  
    mov      rdi, 0       ; exit code 0  
    syscall            ; invoke operating system to exit
```

GLOBAL _start

Why should this now work?



Local Variables

- Variables declared in procedure → allocated on stack

```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo
```

```
→ int main( int argc, char** argv) {  
    foo();  
} // EndMain
```


(info on
call to
main)

Higher addresses

Local Variables

- Variables declared in procedure → allocated on stack


```
→ void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
→ int main( int argc, char** argv) {  
    foo();  
} // EndMain
```



Local Variables


- Variables declared in procedure → allocated on stack

```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv) {  
    foo();  
} // EndMain
```




Local Variables

- Variables declared in procedure → allocated on stack



```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv ) {  
    foo();  
} // EndMain
```




Base Pointer & Stack Frame

- Each function call: its own area on stack
 - Known as a "**stack frame**"
- Local variables create a problem for **rsp**
 - Additional data on top of current return address
 - **rsp** (top of stack) not always pointing to return address
- We need a second register = **rbp** "**base pointer**"
 - A.k.a frame pointer
 - Return address: just before where **rbp** points on the stack
 - Used to access local variables + debugging
- **rbp** needs to be saved after **call** / restored before **ret**
 - Responsibility of callee function


Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack




Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack




Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack



Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack





Base Pointer & Stack Frame


- Variables declared in procedure → allocated on stack



```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv ) {  
    foo();  
} // EndMain
```



rsp →





Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack



```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv ) {  
    foo();  
} // EndMain
```






Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack



```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv ) {  
    foo();  
} // EndMain
```





Base Pointer & Stack Frame

- Variables declared in procedure → allocated on stack



```
void foo(void) {  
    char my_string[] = "Hello World!\n";  
    printf(my_string);  
} // End foo  
  
int main( int argc, char** argv ) {  
    foo();  
} // EndMain
```



(info on
call to
main)



old rbp

"Hello
world!\n"+0x
0



old rbp

local var of
printf + ...

rbp

rsp

Saving / Restoring rbp

- Uses **push** & **pop**
- Function **prologue** (just after **call**)

→ **push** rbp
 mov rbp, rsp
 sub rsp, <space for local variables> ;

- Function **epilogue** (just before **ret**)

→ **mov** rsp, rbp
 pop rbp
 ret r

- On x86 > 8186 compound op : **enter** / **leave**
 - **enter** not used, too slow

Saving Registers

- When issuing a `call foo` code jumps to location 'foo'
 - 'foo' likely to use **registers**
 - **any register value** before `call` **might be overwritten**
 - need to **save registers** that must be preserved
 - Using `push` and `pop` (e.g. `push rax`)

- Who should save / restore what
 - The work can be done on the caller's or the callee's side
 - Or shared
 - Who does what is defined by **calling conventions**
(also covers order of parameters on stack, etc., see later)

Passing Parameters


- 2 mechanisms to transmit parameters (“*transportation*”)
 - Through some **registers** (default), or
 - Through the **stack** (when not enough regs)
- 2 kinds of parameters (“*change propagation*”)
 - By **value** (changes not propagated back to caller)
 - Actual value directly passed in register or on stack
 - By **reference** (aka address) (change propagated back)
 - Register or stack contains an address
 - Actual value obtained after **indirection**
(means 2 indirections if stack frame is used)
- $2 \times 2 = 4$ possibilities

By Value vs. By Reference

- Parameter passed by **value**
 - Value of parameter **copied** into reg **or** onto the stack
 - Change to parameter not visible on return
- Parameter passed by **reference**
 - Needs to pass an address (= a pointer in C)
 - Access to value uses indirection (“dereference”)
 - Changes to parameter will be visible on return
- Note: in Java
 - Everything is a reference, except basic types (values)


What kind of passing is this?

```
SECTION    .data
y:  db   65
...
SECTION    .text
foo:
    add    ax,2          ; x += 2
    ret
_start:
    mov    ax,[y]
    call   foo
```



What kind of passing is this?

```
SECTION    .data
y:  db   65
...
SECTION    .text
foo:
    add    [rax], BYTE 2 ; x += 2
    ret
_start:
    mov    rax,y
    call   foo
```



Returning Value

- Most frequent approach
 - using registers (**rax**, etc.)
- Alternative: on the stack
 - Reserve some space for the return value before calling

Parameters in C (Linux)

- Optimized approach : Using specific registers
 - Input in : rdi, rsi, rdx, rcx, r8, r9
 - Output in : rax (+ rdx if needed)
 - Stack used if parameters / return values do not fit
- Precise rules on how this works
 - Registers need to be saved / restored by callee or caller
 - Depends on language / ISA
 - Linux : specified in **ABI : Application Binary Interface**
 - rbx, rsp, rbp = callee-saved registers
 - The rest : “scratch”, might be overwritten (“clobbered”)

Passing Parameters by Register

■ Example: Decompiling foo() from first example

00000000004004e4 <foo>:

```
void foo(void) {
    4004e4:    55                      push    rbp
    4004e5:    48 89 e5                mov     rbp, rsp
    printf("Executing foo\n");
    4004e8:    bf 4c 06 40 00        mov     edi, 0x40064c
    4004ed:    e8 ee fe ff ff        call    4003e0 <puts@plt>
    printf("Returning from foo\n");
    4004f2:    bf 5a 06 40 00        mov     edi, 0x40065a
    4004f7:    e8 e4 fe ff ff        call    4003e0 <puts@plt>
} // End foo
4004fc:    c9                      leave
4004fd:    c3                      ret
```

→ edi is used to pass the @ of the string to be printed


Summary

As the end of this session you should:

- Be able to explain the role of the stack
- Understand the role of `call` and `ret`
- Be able to understand and manipulate a stack with
`rsp`, `pop`, `push`
- Be able to describe and analyse the structure of stack frames, and the role of `rbp`
- Explain how local variables are allocated in a frame
- Be able to explain how parameters are passed and returned (stack vs. registers, value vs. references)

Bonus Material

- Not exam material




Using Stack to Pass Parameters

- Parameters pushed before the **call** (w/ **push**)
 - Parameter area liberated on return
 - Either with **ret <size to liberate>** (callee clean up)
 - Or by caller (caller clean up, C)

- Parameters and local values **accessed** using **rbp + offset**
 - "based" or "indexed" addressing mode: **rbp[offset]**
 - Positive offsets: parameters
 - Negative offsets: local variable


Stack Layout, Stack-based Passing

Stack growing
from higher to
lower addresses




bar's stack frame

foo's stack frame



Parameters
passed on
stack, no
return value

Who is calling who?



49

Returning Value On the Stack

■ Stack Layout

